


[Apache](#) > [Hadoop](#) > [Core](#) > [common](#) > [docs](#) > [r0.20.203.0](#)



Last Published: 05/04/2011 17:00:30



[PDF](#)

Single Node Setup

- [Purpose](#)
- [Prerequisites](#)
 - [Supported Platforms](#)
 - [Required Software](#)
 - [Installing Software](#)
- [Download](#)
- [Prepare to Start the Hadoop Cluster](#)
- [Standalone Operation](#)
- [Pseudo-Distributed Operation](#)
 - [Configuration](#)
 - [Setup passphraseless ssh](#)
 - [Execution](#)
- [Fully-Distributed Operation](#)

Purpose

This document describes how to set up and configure a single-node Hadoop installation so that you can quickly perform simple operations using Hadoop MapReduce and the Hadoop Distributed File System (HDFS).

Prerequisites

Supported Platforms

- GNU/Linux is supported as a development and production platform. Hadoop has been demonstrated on GNU/Linux clusters with 2000 nodes.
- Win32 is supported as a *development platform*. Distributed operation has not been well tested on

Win32, so it is not supported as a *production platform*.

Required Software

Required software for Linux and Windows include:

1. Java™ 1.6.x, preferably from Sun, must be installed.
2. **ssh** must be installed and **sshd** must be running to use the Hadoop scripts that manage remote Hadoop daemons.

Additional requirements for Windows include:

1. [Cygwin](#) - Required for shell support in addition to the required software above.

Installing Software

If your cluster doesn't have the requisite software you will need to install it.

For example on Ubuntu Linux:

```
$ sudo apt-get install ssh
$ sudo apt-get install rsync
```

On Windows, if you did not install the required software when you installed cygwin, start the cygwin installer and select the packages:

- openssh - the *Net* category

Download

To get a Hadoop distribution, download a recent [stable release](#) from one of the Apache Download Mirrors.

Prepare to Start the Hadoop Cluster

Unpack the downloaded Hadoop distribution. In the distribution, edit the file `conf/hadoop-env.sh` to define at least `JAVA_HOME` to be the root of your Java installation.

Try the following command:

```
$ bin/hadoop
```

This will display the usage documentation for the **hadoop** script.

Now you are ready to start your Hadoop cluster in one of the three supported modes:

- Local (Standalone) Mode
- Pseudo-Distributed Mode
- Fully-Distributed Mode

Standalone Operation

By default, Hadoop is configured to run in a non-distributed mode, as a single Java process. This is useful for debugging.

The following example copies the unpacked `conf` directory to use as input and then finds and displays every match of the given regular expression. Output is written to the given `output` directory.

```
$ mkdir input
$ cp conf/*.xml input
$ bin/hadoop jar hadoop-examples-*.jar grep input output
'dfs[a-z.]+'
$ cat output/*
```

Pseudo-Distributed Operation

Hadoop can also be run on a single-node in a pseudo-distributed mode where each Hadoop daemon runs in a separate Java process.

Configuration

Use the following:

`conf/core-site.xml`:

```
<configuration>
  <property>
    <name>fs.default.name</name>
    <value>hdfs://localhost:9000</value>
  </property>
</configuration>
```

`conf/hdfs-site.xml`:

```
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
</configuration>
```

`conf/mapred-site.xml`:

```
<configuration>
  <property>
    <name>mapred.job.tracker</name>
    <value>localhost:9001</value>
  </property>
</configuration>
```

```
</property>
</configuration>
```

Setup passphraseless ssh

Now check that you can ssh to the localhost without a passphrase:

```
$ ssh localhost
```

If you cannot ssh to localhost without a passphrase, execute the following commands:

```
$ ssh-keygen -t dsa -P '' -f ~/.ssh/id_dsa
$ cat ~/.ssh/id_dsa.pub >> ~/.ssh/authorized_keys
```

Execution

Format a new distributed-filesystem:

```
$ bin/hadoop namenode -format
```

Start the hadoop daemons:

```
$ bin/start-all.sh
```

The hadoop daemon log output is written to the `${HADOOP_LOG_DIR}` directory (defaults to `${HADOOP_HOME}/logs`).

Browse the web interface for the NameNode and the JobTracker; by default they are available at:

- NameNode - <http://localhost:50070/>
- JobTracker - <http://localhost:50030/>

Copy the input files into the distributed filesystem:

```
$ bin/hadoop fs -put conf input
```

Run some of the examples provided:

```
$ bin/hadoop jar hadoop-examples-*.jar grep input output
'dfs[a-z.]+'
```

Examine the output files:

Copy the output files from the distributed filesystem to the local filesystem and examine them:

```
$ bin/hadoop fs -get output output
$ cat output/*
```

or

View the output files on the distributed filesystem:

```
$ bin/hadoop fs -cat output/*
```

When you're done, stop the daemons with:

```
$ bin/stop-all.sh
```

Fully-Distributed Operation

For information on setting up fully-distributed, non-trivial clusters see [Cluster Setup](#).

Java and JNI are trademarks or registered trademarks of Sun Microsystems, Inc. in the United States and other countries.

Last Published: 05/04/2011 17:00:30

Copyright © 2008 [The Apache Software Foundation](#).